

Introduction to High-Performance Computing (HPC)


HPC cluster

“High Performance Computing most generally refers to the practice of aggregating computing power in a way that delivers much higher performance than one could get out of a typical desktop computer or workstation in order to solve large problems in science, engineering, or business.”

Need for HPC cluster?

- **Resources:** Your computer does not have the resources to run the desired NGS analysis
- **Speed:** You want to produce results faster than your computer can
- **Software:** You need software that is unavailable or unusable on your computer

HPC cluster



Need for HPC cluster?

- **Resources:** Your computer does not have the resources to run the desired NGS analysis
- **Speed:** You want to produce results faster than your computer can
- **Software:** You need software that is unavailable or unusable on your computer

HPC cluster components

Nodes: Individual computers in the cluster

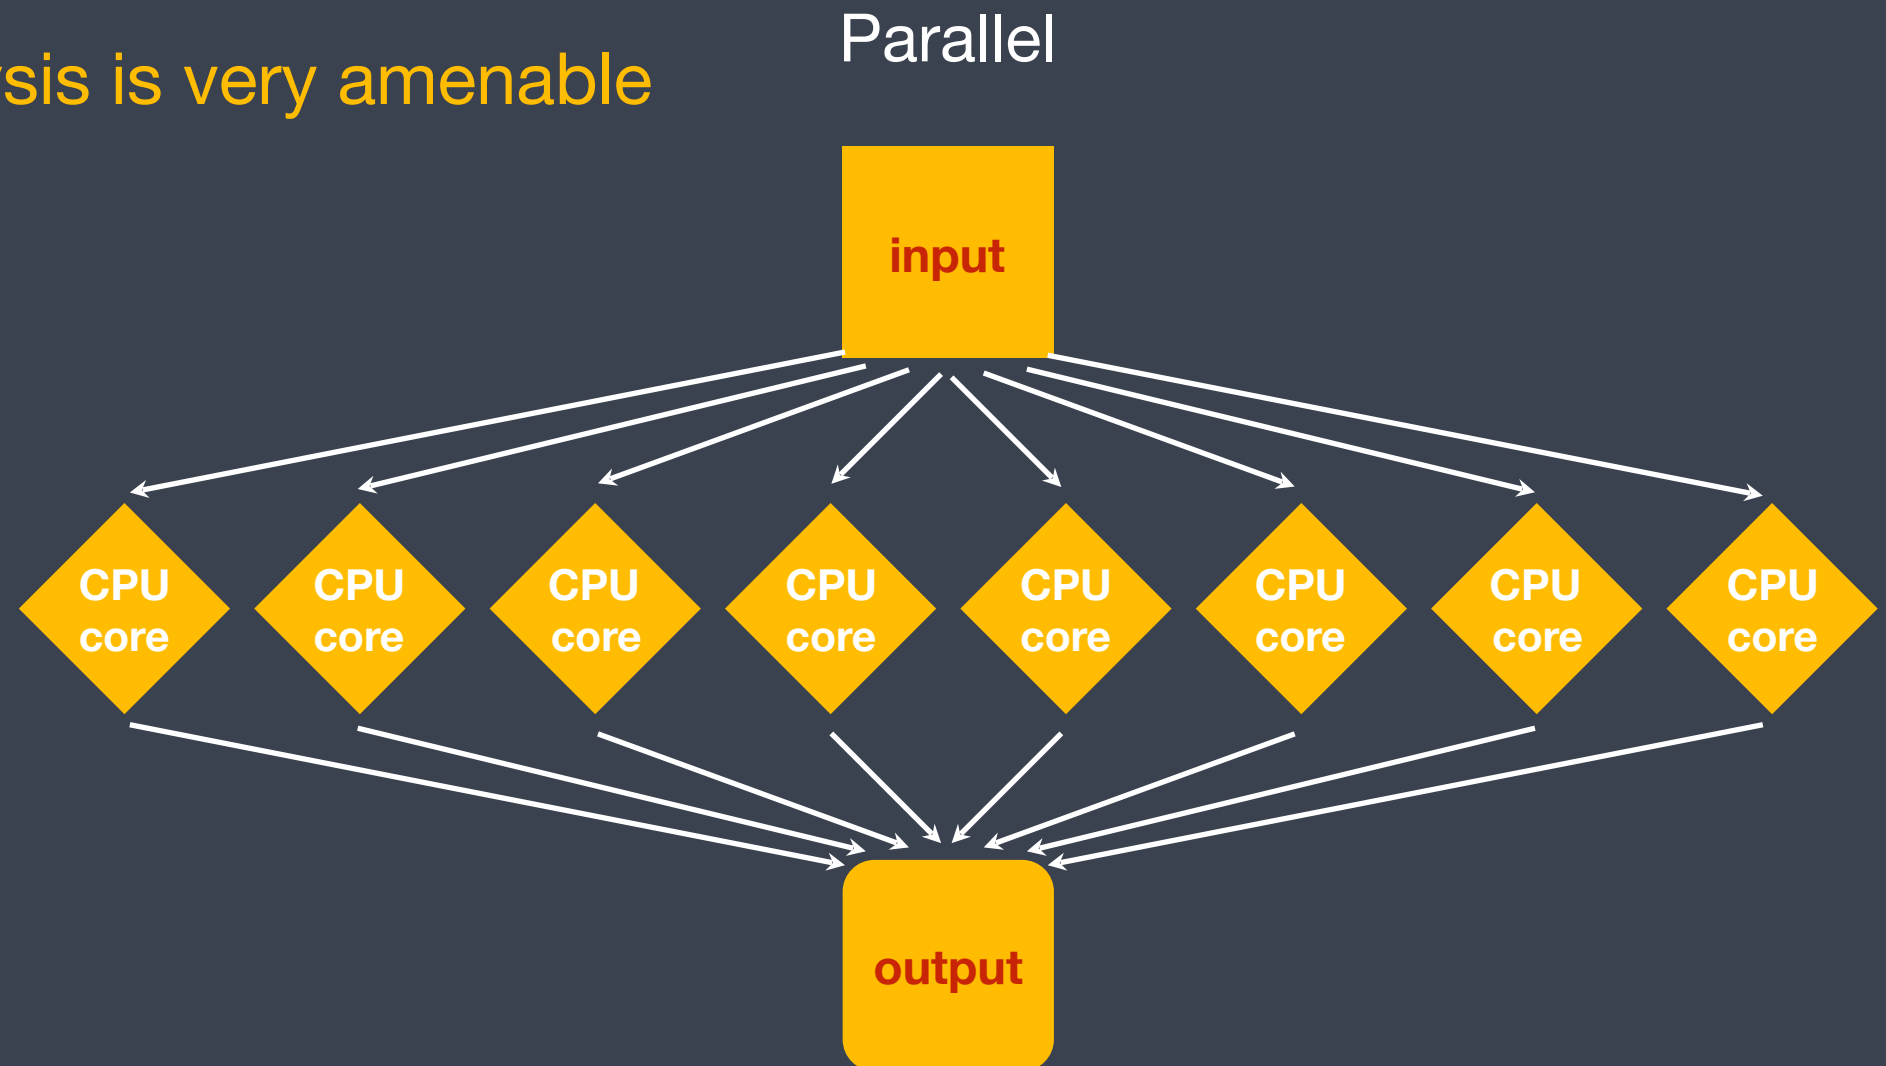
Cores: individual processing units available within each CPU of each Node

e.g. a “Node” with eight “quad”-core CPUs = 32 cores for that node.

Shared disk: storage that can be shared (and accessed) by all nodes

Multi-threading

NGS data analysis is very amenable to this strategy.



Faster and more efficient

HPC cluster

- multi-user, shared resource
- lots of nodes = lots of processing capacity + lots of memory
- a system like this requires constant maintenance and upkeep.



Please see our [Orchestra status](#) page for known issues.

The **Orchestra** platform provides UNIX-based high performance computing, web hosting, and database hosting services at Harvard Medical School.

Orchestra and its associated services are managed by the [Research Computing Group](#), part of the HMS [Information Technology Department](#).

Please [submit a request](#) via the RC web site for help with Orchestra or feedback. You can also subscribe to the [mailing list](#) for Orchestra users.

Wiki page: <https://wiki.med.harvard.edu/Orchestra>

Introduction to High Performance Computing and Orchestra

HMS Research Computing
Spring 2016



What is Orchestra?

- Tech specs:
 - Over 550 compute nodes
 - Over 7,500 cores
 - Over 42TB RAM
- CentOS 6 Linux
- LSF scheduler
- Total 28+PB storage



Using Orchestra!

1. Logging in to remote machines (securely)

- When logging in we used the “ssh” command,
ssh stands for Secure Shell
- **ssh** is a protocol for data transfer that is secure, i.e the data is encrypted as it travels between your computer and the cluster (remote computer)
- Commonly used commands that use the **ssh** protocol for data transfer are, **scp** and **sftp**

2. Using & installing software

Using Software Environment Modules

- Most “software” on Orchestra is installed as an environment module.
- An environment module makes the necessary modifications in \$PATH to make sure the program runs without any issues.
- Allows for clean, easy loading, including most dependencies, and switching versions.

Using Software Environment Modules

```
$ module avail
```

```
$ module avail seq/
```

```
$ module avail stats/
```

```
$ module avail seq/bowtie/
```

What software
environment modules are
available on Orchestra?

- Loading modules

```
$ module load seq/bowtie/2.0.6
```

- See all modules that have been loaded

```
$ module list
```

- Unloading modules

```
$ module unload seq/bowtie/2.0.6
```

3. The Job Scheduler, LSF

LSF: Fair Sharing

- **Load Sharing Facility:** distributes jobs across Orchestra fairly
- Ensures that no single user or core monopolizes Orchestra
- Users are assigned dynamic priorities
- Queues also have priorities
- Submitting lots of jobs reduces your fairshare priority
- Even if many jobs are pending, your jobs will start quickly provided you have not submitted many jobs

Choosing the proper resources for your job with
the appropriate `bsub` options

The “bsub”

```
$ bsub -q queue -W hr:min job/command
```

- ♦ Necessary to specify:
 - q (queue)
 - W (runtime in hr:min)

bsub options

`-W 3:00` *#time requested*
`-q multicore` *#name of the queue*
`-n 4` *#number of cores requested*
`-R "rusage[mem=8000]"` *#memory requested in MB*
`-J jobname`
`-a openmpi` *#type of mpi run*
`-Is` *#interactive shell*
`-e %J.err` *#send errors to file*
`-o %J.out` *#send screen output to file*
`-N` *#notify when job completes*

Runtime Limit ($b_{sub} - W$)

- In hours:minutes
- Runtimes are subject to the maximum time permitted per queue (see table)
- If your job exceeds your runtime, your job will be killed ☹️
- Running many jobs that finish quickly (less than a few minutes) is suboptimal and may result in job suspension, contact RC to learn how to batch jobs

Multithreading (`bsub -n`)

- A single CPU can execute multiple processes (threads) concurrently
- “-n” indicates how many cores are requested
- Jobs that are overefficient (use more cores than reserved) jeopardize the health of a node
- Reserve the same amount of cores in your job and your `bsub`!

CPU Limit

- Amount of seconds the cluster works on your job (calculated by LSF)
- Ncores * Runlimit (-n * -W)
- Common error:

```
bsub -q short -W 8:00 tophat -p 8
```

tophat asks for 8 cores (-p 8) but only 1 core requested (no -n), job killed in 1 hour (8/8)

Reserving Memory (`bsub -R`)

- Most nodes have 90GB memory available over all cores, some have more
- Make a resource request with
 - R “`rusage[mem=8000]`” (memory requested in MB)
- Memory multiplies by cores requested, so
 - n 4 -R “`rusage[mem=16000]`” reserves 64GB memory
- Asking for more memory may cause jobs to pend longer
- `TERM_MEMLIMIT` errors indicate that not enough memory was reserved

Shared Queues (`bsub -q`)

- ***mpi*** queue if you have an MPI parallel job
- ***priority*** queue if you have just one or two jobs to run
- ***mcore*** queue if you have multi-core jobs to run.
- ***short*** queue if your jobs will take less than 12 hours to run.
- Else: ***long*** queue.

<https://wiki.med.harvard.edu/Orchestra/ChoosingAQueue>

Job submission scripts

Shell Scripts: The Basics

```
#!/bin/sh                                #always at the top

#commands with options
tophat -p 4 -o ./mytophatdir1 hg19 file1_1.fastq file1_2.fastq
tophat -p 4 -o ./mytophatdir2 hg19 file2_1.fastq file2_2.fastq
```

Save as `myshellscript.sh`

Run this script in an interactive session as `sh myshellscript.sh`

OR

Submit it as a job to the LSF scheduler

```
$ bsub -q mcore -W 12:00 -n 4 -N -e %J.err -o %J.out sh myshellscript.sh
```

Creating a Job Submission Script

```
#!/bin/sh
#BSUB -q mcore
#BSUB -W 12:00
#BSUB -o %J.out
#BSUB -e %J.err
#BSUB -N
#BSUB -n 4
#BSUB -R "rusage[mem=12000]"

module load seq/tophat2.1.1

tophat -p 4 -o ./mytophatdir1 hg19 file1_1.fastq file1_2.fastq
```

Save as **myshellscript.lsf**

Run as **\$ bsub < myshellscript.lsf**

Orchestra will notify you when the job is done, or if there is an error

Managing jobs and getting information about
submitted/running jobs

Monitoring Jobs

- List info about jobs/their status:

```
mfk8$loge:~$ bjobs
```

-r (running jobs)

-p (pending jobs)

-l (command entered, long form)

Terminating Jobs

- Terminate a job (jobid given at submission)


```
mfk8$loge:~$ bkill jobid
```

- Terminate all jobs (be careful!)

```
mfk8$loge:~$ bkill 0
```

4. Filesystems and storage


Filesystems and storage



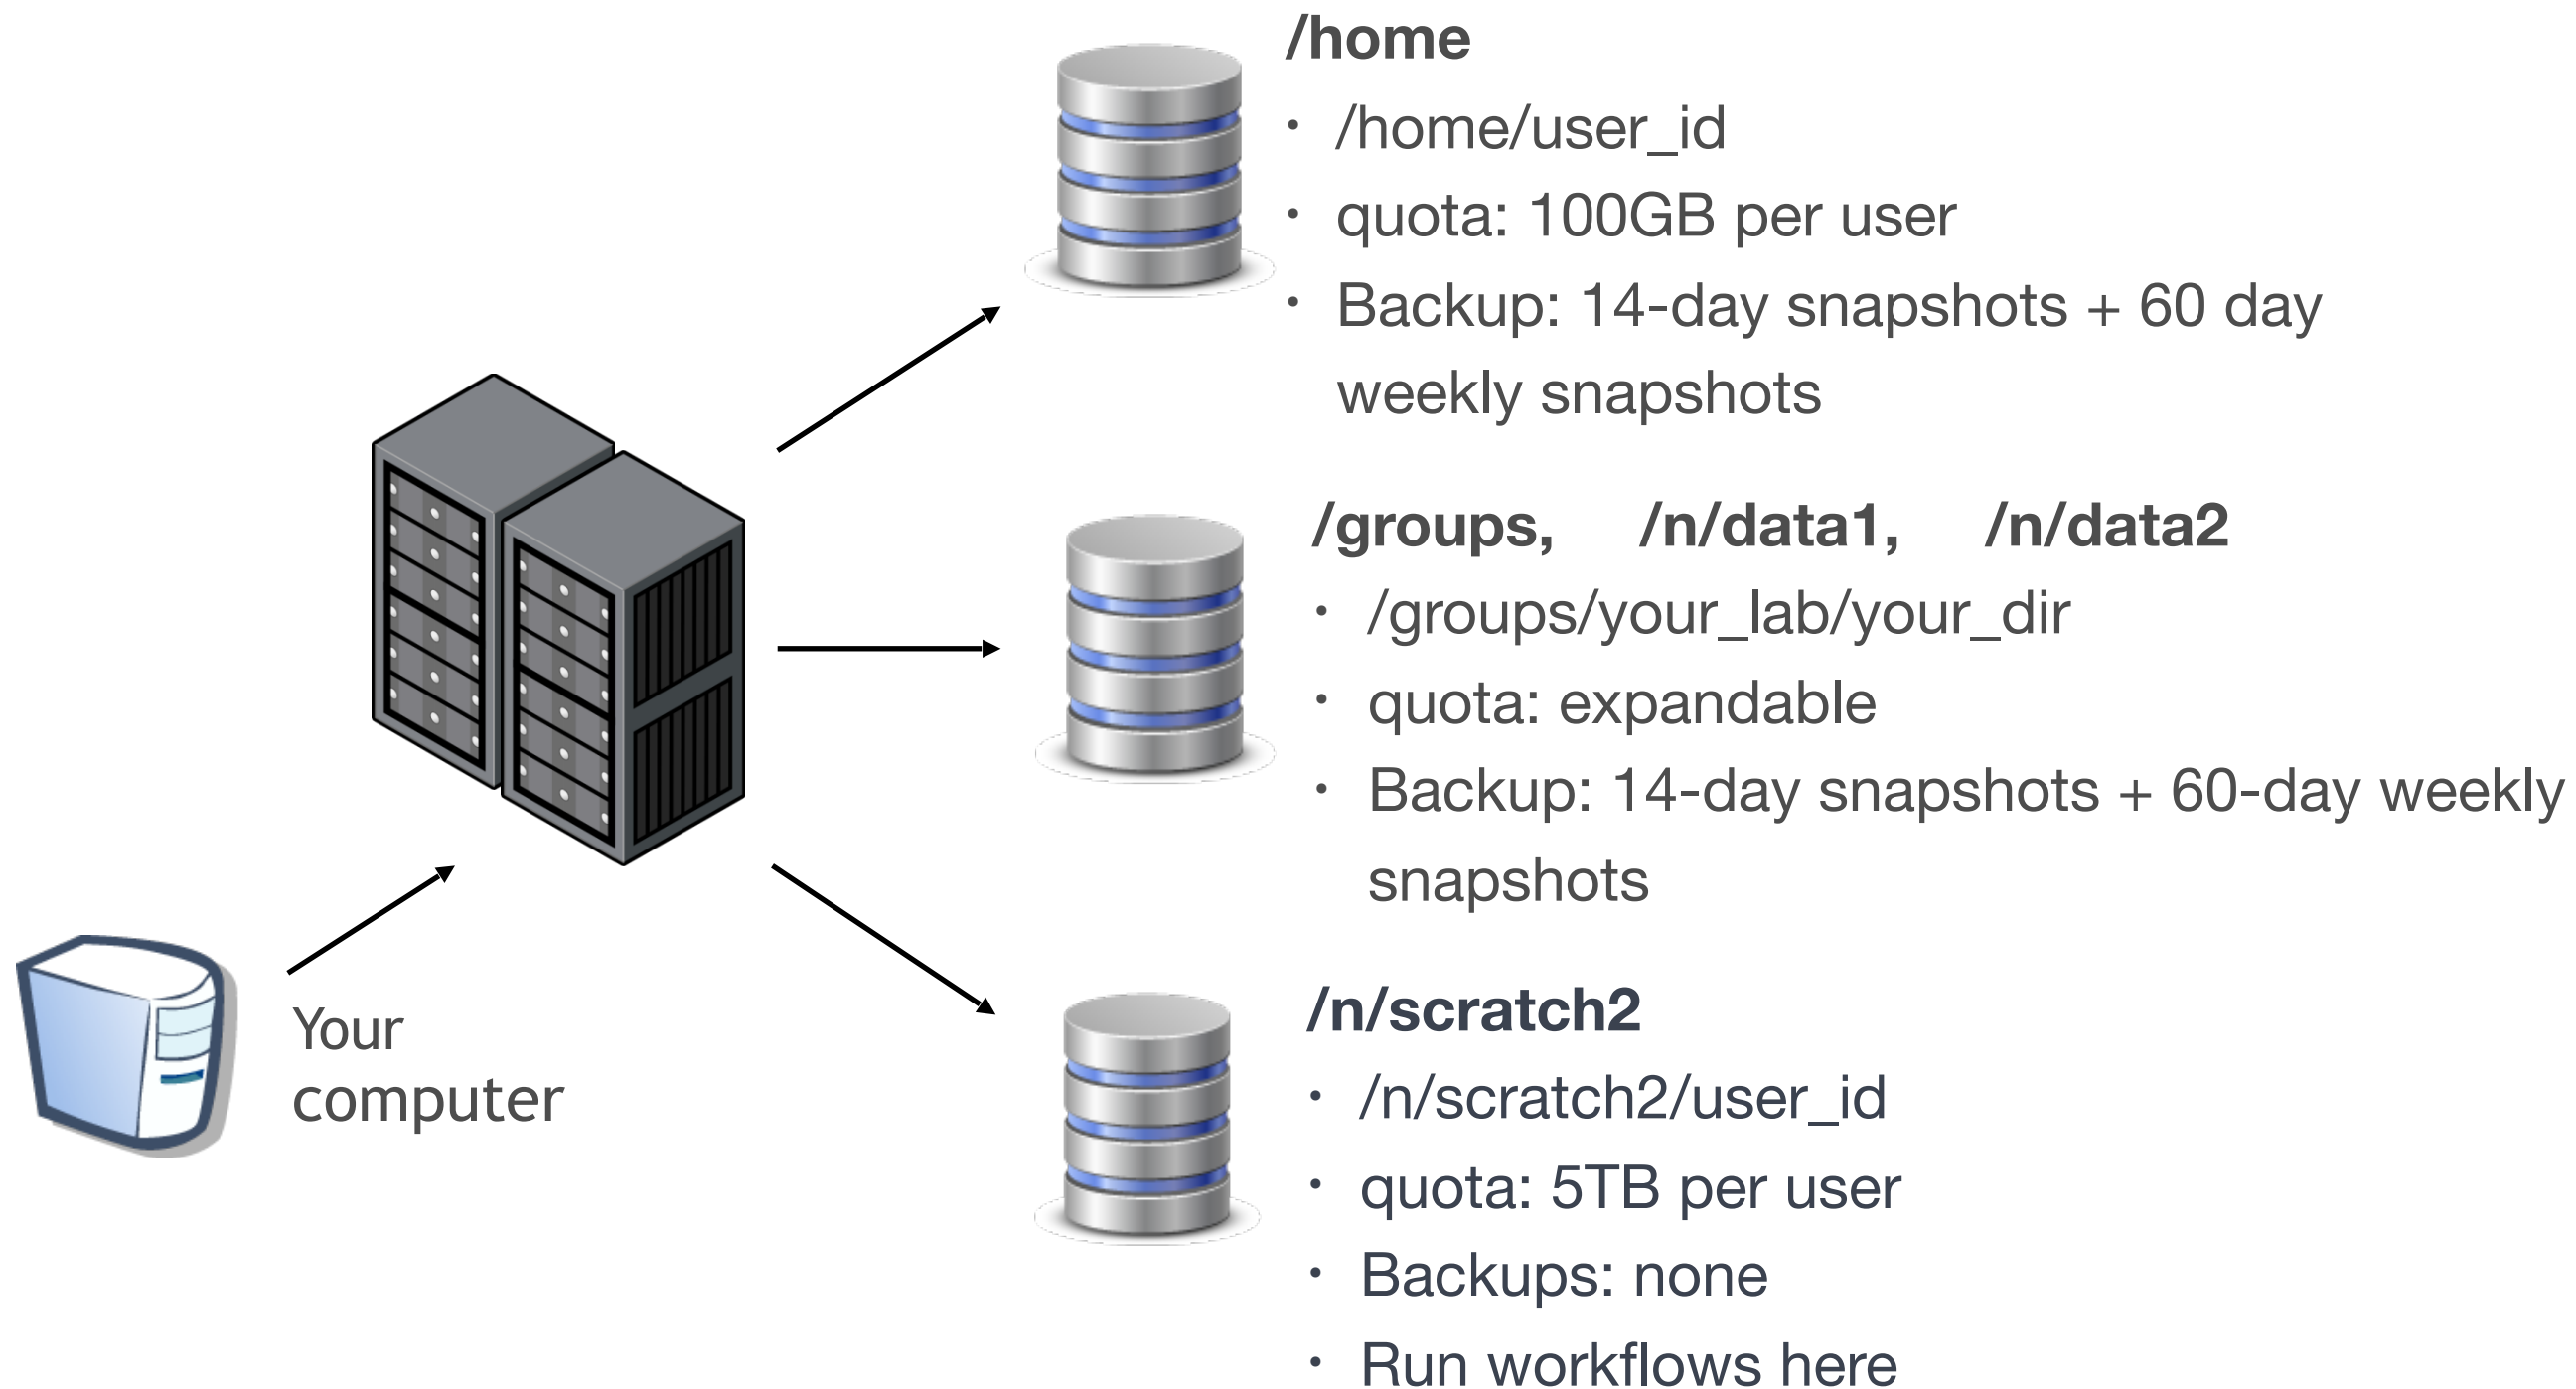
Filesystems and storage

- Storage on HPC systems is organized differently than on your personal machine
- Physical disks are bundled together into a virtual volume; this volume may represent a single filesystem, or may be divided up, or partitioned, into multiple filesystems
- Filesystems are accessed over the internal network

Filesystems and storage



Storage on Orchestra



Checking storage usage

- To check your storage available:

`$ quota`

- Home directory (/home/username): you get 100GB.
- Group directories: space varies, can increase.

`/groups/groupname`

`/n/data1`

`/n/data2`

For more direction:

- <https://wiki.med.harvard.edu/Orchestra/NewUserGuide>
- <http://rc.hms.harvard.edu>
- rchelp@hms.harvard.edu
- Office Hours: Wednesdays 1-3p Gordon Hall 500